Credal Prediction for Large Language Models

Liam Heppner

March 22, 2026

Department of Mathematics,
Informatics and Statistics
Institute of Informatics

Artificial Intelligence and
Machine Learning

Bachelor's Thesis

Credal Prediction for Large Language Models

Liam Heppner

Reviewer Prof. Dr. Eyke Hüllermeier
Institute of Informatics
LMU Munich

Supervisor Paul Hofman

March 22, 2026

Liam Heppner
Credal Prediction for Large Language Models
Bachelor's Thesis, March 22, 2026
Reviewers: Prof. Dr. Eyke Hüllermeier
Supervisor: Paul Hofman

LMU Munich
Department of Mathematics, Informatics and Statistics
Institute of Informatics
Artificial Intelligence and Machine Learning (AIML)
Akademiestraße 7
80799 Munich

Abstract

Standard uncertainty estimation in Large Language Models (LLMs) relies on single, fixed probability distributions. This "point-estimate" approach often conflates benign linguistic diversity with factual ignorance, failing to reliably quantify uncertainty. In this thesis, we introduce *Credal Semantic Entropy*, a framework grounded in Imprecise Probabilities. We propose a method to construct credal sets—convex sets of probability distributions—based on the statistical notion of relative likelihood. Instead of relying on a single temperature, we identify a set of temperatures that maximize the likelihood of calibration data. We cluster generated sequences by semantic meaning and derive interval-valued probabilities, computing the Upper Semantic Entropy to capture a worst-case estimate of uncertainty. We empirically evaluate our approach on the CoQA benchmark using the OPT-6.7B model. Experiments reveal that the credal framework is highly sensitive to relative likelihood thresholds (α-cuts). Under our tested configurations, single-point baselines outperform Credal Semantic Entropy in uncertainty estimation (measured via AUROC).

Acknowledgement

Large Language Models (LLMs) were used to support several stages of this work. Their assistance included co-coding and troubleshooting during development, helping to brainstorm approaches and refine ideas, and providing guidance for resolving code bugs. LLMs were also used to improve the clarity of written content by correcting spelling, grammar, and phrasing. All substantive decisions, analysis, and final interpretations were made by the author.

Contents

1. Introduction **1**
1.1. Motivation . . . 1
1.2. Problem Statement . . . 1
1.3. Contributions . . . 2
1.4. Thesis Structure . . . 2

2. Background **5**
2.1. Uncertainty in Machine Learning . . . 5
2.1.1. Aleatoric and Epistemic Uncertainty . . . 5
2.1.2. Standard Metrics: Shannon Entropy and Maximum Softmax probability (MSP) . . . 6
2.2. Imprecise Probabilities and Credal Sets . . . 7
2.2.1. Introduction to Credal Sets . . . 7
2.2.2. The Geometry of Credal Sets . . . 7
2.2.3. Computing Uncertainty on Credal Sets . . . 8
2.3. Large Language Models and Generation . . . 9
2.3.1. Transformers and Next-token prediction . . . 9
2.3.2. Decoding Strategies and the Role of Temperature . . . 10
2.4. Quantifying Uncertainty in LLMs . . . 12
2.4.1. Negative Log-Likelihood in LLMs . . . 12
2.5. Relative Likelihood and Valid Temperature Sets . . . 13
2.5.1. Relative Likelihood . . . 13
2.5.2. Constructing α-Cuts for Temperatures . . . 14

3. Related Work **15**
3.1. Probability-Based Methods . . . 15
3.2. Lexical Similarity . . . 16
3.3. Self-Evaluation . . . 16
3.4. Semantic Entropy . . . 16
3.4.1. Semantic Clustering . . . 17
3.4.2. Computing the Semantic Entropy . . . 17
3.5. Credal Prediction Based on Relative Likelihood . . . 18
3.5.1. Relative Likelihood Ensembles . . . 19
3.5.2. From Ensembles to LLM Temperatures . . . 19

4. Methodology: Credal Semantic Entropy via Temperature Scaling **21**
4.1. Overview of the Pipeline . . . 21
4.2. Valid Temperature Selection (the $\alpha-$Cut) . . . 22
4.2.1. Computing the Average Negative Log-Likelihood . . . 22
4.2.2. Deriving Relative Likelihoods . . . 23

4.2.3. The α-Cut . . . 23
4.3. Generation and Semantic Clustering . . . 24
4.3.1. Multi-Temperature Sequence Generation . . . 24
4.3.2. Bidirectional Entailment and Equivalence Classes . . . 24
4.3.3. Constructing the Cluster Probability Distributions . . . 25
4.3.4. Quantifying Credal Semantic Uncertainty . . . 26

5. Experiments **29**
5.1. Experimental Setup . . . 29
5.1.1. Dataset: CoQA . . . 29
5.1.2. Model Architecture . . . 29
5.1.3. Baselines . . . 30
5.1.4. Evaluation Metric: AUROC . . . 30
5.2. Results . . . 31
5.2.1. Evaluation: 20-Step Interval . . . 31
5.2.2. Evaluation: 50-Step Interval . . . 32
5.2.3. Additional Evaluation: Local Analysis at $\tau = 0.5$ . . . 33
5.3. Discussion . . . 34

6. Conclusion **37**
6.1. Summary . . . 37
6.2. Limitations . . . 38
6.3. Future Work . . . 38
6.3.1. Generalization Across Domains and Architecture . . . 38
6.3.2. Multi-Dimensional Decoding Ensembles . . . 39
6.3.3. Methodological Extensions: Ensembles and Black-Box Settings 39

A. Appendix **41**
A.1. Experimental Details . . . 41

Bibliography **43**

List of Figures **49**

List of Tables **51**

1 Introduction

1.1 Motivation

Large Language Models (LLMs) have fundamentally transformed the landscape of Artificial Intelligence, demonstrating unprecedented capabilities in natural language understanding, reasoning and generation [Bro+20; Wei+23]. LLMs are increasingly deployed in critical domains ranging from healthcare [Bed+25] to legal analysis [Sii+25]. However, despite their fluency, these models suffer from a critical reliability failure: the tendency to generate plausible-sounding but factually incorrect statements, a phenomenon widely known as hallucinations [Hua+23].

For high-stakes applications, accuracy alone is insufficient. A system must also be self-aware of its limitations [Bom+22]. Unfortunately, modern LLMs are frequently "confidently wrong", assigning high probability scores to their hallucinations due to poor calibration [Guo+17; DD20] and the nature of their training objectives [Ji+23].

Standard uncertainty metrics in Nautural Language Generation (NLG), such as predictive entropy [MG21] or lexical similarity [Fom+20], often fail to detect these errors because they rely on a single, precise probability distribution. This "point-estimate" approach mathematically conflates two distinct sources of uncertainty [HW21]: *aleatoric uncertainty* (the natural diversity of valid human language) and *epistemic uncertainty* (the model's lack of knowledge). When an LLM distributes its probability mass across five different ways of saying the correct answer for a given prompt, standard metrics flag this as high uncertainty, confusing benign linguistic flexibility with factual ignorance. To overcome this limitation, we build upon the semantic clustering framework introduced by Kuhn et al. [KGF23]. By shifting the unity of measurement from exact lexical strings to underlying meanings, we propose a novel method that establishes probability intervals over these semantic clusters to capture uncertainty for a conversational question-answering task, specifically, CoQA [RCM19].

1.2 Problem Statement

The core problem addressed in this thesis is the unreliability of single-point probability estimates for uncertainty quantification (UQ) in free-form question-answering.

Current state-of-the-art methods, such as Semantic Entropy [KGF23], attempt to mitigate this by grouping lexically diverse answers into semantic clusters. However, even Semantic Entropy relies on a single, static probability distribution generated at a fixed temperature. If this underlying distribution is miscalculated, i.e., peaked around an incorrect fact, the metric will fail to signal a warning. To reliably quantify uncertainty, we need a framework that moves beyond precise probabilities and considers a set of plausible model states.

1.3 Contributions

This thesis proposes a novel methodology for quantifying epistemic uncertainty by bridging the gap between imprecise probabilities and LLM decoding.

Drawing on the theory of *Credal Prediction based on Relative Likelihood* [Löh+25], we introduce a method to construct a credal set of semantic distributions. Our specific contributions are as follows:

- **Valid Temperature Selection via Relative Likelihood**: We propose a method to identify a valid range of decoding temperatures. By utilizing relative likelihood and applying an α-cut, we systematically discard implausible distributions with a relative likelihood below α. We treat the α-threshold as a hyperparameter, increasing the α-threshold narrows the valid temperature range, whereas decreasing it broadens the range.

- **Construction of Semantic Credal Sets**: We extend the concept of Semantic Entropy to operate over Interval-Valued Probabilities. By computing the minimum and maximum probability of semantic clusters across valid temperatures, we construct a convex credal set that captures the model's true epistemic ambiguity.

- **Empirical Validation**: We evaluate our framework on the CoQA benchmark using the OPT-6.7B model. Through experiments across multi-step temperature-grids, we analyze the sensitivity of the α-cut threshold and compare the uncertainty estimation performance of Credal Semantic Entropy against established single-point baselines.

1.4 Thesis Structure

The remainder of this thesis is organized as follows:

- **Chapter 2 Background**: establishes the theoretical foundations necessary for this work. It reviews the architecture of LLMs, explains autoregressive decoding strategies, and introduces the mathematical framework of imprecise probabilities and credal sets.

- **Chapter 3 Related Work**: surveys the current landscape of UQ. It contrasts traditional point-estimate methods with recent advancements in semantic consistency, specifically focusing on Semantic Entropy as the primary baseline for our research.

- **Chapter 4 Methodology**: formalizes our novel contribution: *Credal Semantic Entropy*. It details the algorithm for selecting valid temperatures via relative likelihood and deriving uncertainty bounds from interval-valued probability distributions.

- **Chapter 5 Experiments**: presents the experimental setup and results, analyzing the performance of our method against strong baselines.

- **Chapter 6 Conclusion**: summarizes the core empirical findings of this thesis, discusses the practical and theoretical limitations of the proposed credal framework and outlines promising directions for future research.

2 Background

This chapter establishes the theoretical and methodological foundations necessary to understand the contributions of this thesis. Our work lies at the intersection of Natural Language Processing (NLP) and Uncertainty Quantification (UQ). We first establish the mathematical principles of uncertainty in standard machine learning, followed by the probabilistic mechanisms of modern language models and finally how these domains unite in our methodology.

2.1 Uncertainty in Machine Learning

As LLMs are increasingly deployed in real-world and high-stakes applications, for example in healthcare [Bed+25], their reliability becomes a critical concern. A fundamental challenge with deep neural networks, including LLMs, is that they are prone to overconfidence. They frequently assign high probabilities to incorrect predictions [Guo+17]. Consequently, evaluating a model's performance requires not only measuring its accuracy but also quantifying its **uncertainty**.

2.1.1 Aleatoric and Epistemic Uncertainty

In the field of machine learning, uncertainty can be categorized into two distinct sources: aleatoric and epistemic [KG17; HW21].

Aleatoric Uncertainty arises from the inherent complexity, noise, or randomness in the underlying data generation process. It is irreducible, meaning providing the model with more training data will not eliminate this uncertainty [KG17; HW21]. In a standard classification task, this might occur when two distinct classes share highly similar features, making perfect separation impossible regardless of the model's capacity.

Epistemic uncertainty, on the other hand, stems from a lack of knowledge or information about the best model. This uncertainty occurs when a model encounters out-of-distribution (OOD) inputs or queries about facts absent from its training data. Unlike aleatoric uncertainty, epistemic uncertainty is reducible. Theoretically it can be minimized by providing the model with more comprehensive training data [KG17; HW21].

2.1.2 Standard Metrics: Shannon Entropy and Maximum Softmax probability (MSP)

To quantify these theoretical concepts, we rely on established mathematical metrics. For a standard classification model predicting a class Y given input $\mathbf{x}$, two fundamental measures of predictive confidence are **Shannon Entropy** [Sha48] and **Maximum Softmax probability (MSP)** [HG18].

Shannon Entropy quantifies the fundamental uncertainty inherent in the entire probability distribution over all possible classes $y \in \mathcal{Y}$. For a discrete random variable Y with a probability distribution $P(Y \mid \mathbf{x})$, the Shannon Entropy H is defined as [Sha48]:

$$H(Y) = -\sum_{y \in \mathcal{Y}} P(y \mid \mathbf{x}) \log P(y \mid \mathbf{x}) \tag{2.1}$$

Entropy provides a scalar summary of the distribution's shape.

- **Low Entropy** ($H \to 0$): The distribution is highly concentrated (peaked) on one or a few tokens, indicating that the model is highly certain about its next prediction.

- **High Entropy** ($H \to \log |\mathcal{V}|$): The distribution is flat and uniform across the vocabulary, indicating maximum uncertainty.

Maximum Softmax Probability (MSP)
While Shannon Entropy evaluates the uncertainty across the entire label space $\mathcal{Y}$, the most ubiquitous measure of point-prediction confidence in deep learning is the Maximum Softmax Probability. For a given input $\mathbf{x}$, MSP simply takes the probability of the most likely predicted class $\hat{y}$ as the model's absolute confidence score:

$$MSP(\mathbf{x}) = \max_{y \in \mathcal{Y}} P(y \mid \mathbf{x}) \tag{2.2}$$

MSP provides a straightforward scalar value ranging from $\frac{1}{|\mathcal{Y}|}$ (complete uncertainty) to 1.0 (absolute certainty).

Limitations of Standard Metrics
Shannon Entropy and MSP are highly effective metrics, provided that the underlying probability distribution P is perfectly calibrated and accurately reflects the true state of the world. However, these metrics inherently assume that the model's output distribution is a single, precise, and entirely correct representation of reality [Guo+17; GR07].

As demonstrated by Guo et al. [Guo+17], this assumption of a "precise probability" often breaks down in modern deep learning models, which frequently suffer from severe overconfidence and miscalibration. Furthermore, standard precise probabilities mathematically struggle to distinguish between inherent data noise (aleatoric) and a fundamental lack of knowledge (epistemic) [HW21]. This failure to capture true epistemic misconfidence necessitates more robust mathematical frameworks, such as imprecise probabilities and credal sets[Wal91].

2.2 Imprecise Probabilities and Credal Sets

Traditional probability theory operates on the assumption that all forms of uncertainty can be adequately represented by a single, precise probability distribution [Wal91]. Under this paradigm, an agent or model is required to assign an exact numerical probability to every possible event, regardless of the amount of evidence available. While computationally convenient, this precise framework struggles to model severe epistemic uncertainty. When a model lacks sufficient training data or encounters an ambiguous query, forcing it to produce a single probability distribution often results in arbitrary or misleadingly confident predictions. Precise probabilities mathematically conflate the lack of evidence (epistemic uncertainty) with known stochastic noise (aleatoric uncertainty), masking the model's true state of ignorance [HW21]. Specifically, a single, flat probability distribution cannot distinguish between an inherently ambiguous input and a completely novel one, as both scenarios yield the same low-confidence prediction.

2.2.1 Introduction to Credal Sets

To address the limitations of precise probabilities, the framework of **imprecise probabilities** replaces the single distribution with a set of plausible distributions [Lev80; Wal91]. A closed, convex set of probability distributions is known as a **credal set** [Coz00], denoted by $\mathcal{K}$.

2.2.2 The Geometry of Credal Sets

A credal set $\mathcal{K}$ is geometrically represented as a closed, convex subset of the standard probability simplex Δ. For a classification task with K possible classes $c_1, c_2, \ldots, c_K$ (or semantic clusters), the standard simplex Δ_K is defined as [Pin88]:

$$\Delta_K = \left\{ P \in \mathbb{R}^K \mid \forall k : P_k \geq 0, \sum_{k=1}^{K} P_k = 1 \right\} \tag{2.3}$$

A probability distribution is a single point within this simplex. In contrast, a credal set represents a region of points.

Interval-Valued Probabilities
One of the most common and practical ways to define this region is through **probability intervals**. Instead of assigning a precise value $P(c_k \mid \mathbf{x})$, we assign a lower bound $\underline{P}(c_k \mid \mathbf{x})$ and an upper bound $\overline{P}(c_k \mid \mathbf{x})$ to each class c_k [CHM11].

The credal set induced by these intervals is the set of all distributions consistent with these bounds:

$$\mathcal{K} = \left\{ P \in \Delta_K \mid \underline{P}(c_k \mid \mathbf{x}) \leq P(c_k \mid \mathbf{x}) \leq \overline{P}(c_k \mid \mathbf{x}), \sum_{k=1}^{K} P(c_k \mid \mathbf{x}) = 1 \right\} \tag{2.4}$$

Geometrically, these linear constraints form a **convex polytope** within the simplex [CHM11].

If the intervals are wide (high epistemic uncertainty), the polytope covers a large volume of the simplex. If the intervals are narrow (low uncertainty), the polytopes shrinks toward a single point. This geometric interpretation is fundamental to understanding how we quantify uncertainty in Chapter 4.

2.2.3 Computing Uncertainty on Credal Sets

To quantify the uncertainty of a model's prediction using a credal set, we must extend standard metrics like Shannon Entropy to operate over sets of distributions. Since we have a set of probability distributions $\mathcal{K}$, we can determine the **lower entropy** and **upper entropy** [AM05] by finding the minimum and maximum entropy values across the set [AKM06; HW21]:

$$H_*(\mathcal{K}) = \min_{P \in \mathcal{K}} H(P), \quad H^*(\mathcal{K}) = \max_{P \in \mathcal{K}} H(P) \tag{2.5}$$

We can use these bounds to quantify total uncertainty $U(\mathcal{K})$, aleatoric uncertainty $AU(\mathcal{K})$ and epistemic uncertainty $EU(\mathcal{K})$ [AKM06; HW21]:

$$U(\mathcal{K}) = AU(\mathcal{K}) + EU(\mathcal{K}) \tag{2.6}$$

$$H^*(\mathcal{K}) = H_*(\mathcal{K}) + (H^*(\mathcal{K}) - H_*(\mathcal{K})) \tag{2.7}$$

Here, upper entropy represents total uncertainty [AM05], lower entropy represents aleatoric uncertainty and the difference between them is used as a measure for epistemic uncertainty.

2.3 Large Language Models and Generation

Having established the framework for quantifying uncertainty, we now turn to the models under investigation. Large Language Models (LLMs) represent a class of deep learning models designed to process, understand, and generate human language. At their core, modern LLMs are built upon the Transformer architecture, utilizing mechanisms such as self-attention to model complex dependencies between words across long sequences [Vas+17; Nav+24]. They demonstrate capabilities in reasoning [Wei+23] and coding [Tia+23]. Their fundamental operation remains probabilistic. They predict the likelihood of the next token in a sequence based on the preceding context [Vas+17; BCB16].

2.3.1 Transformers and Next-token prediction

The primary objective of an autoregressive LLM is to model the probability distribution of a target sequence of tokens $\mathbf{y} = (y_1, y_2, \ldots, y_T)$ given an input context or prompt $\mathbf{x}$. The model decomposes the joint probability of the sequence into a product of conditional probabilities using the chain rule of probability [BCB16]:

$$P(\mathbf{y}|\mathbf{x}) = \prod_{t=1}^{T} P(y_t \mid \mathbf{y}_{<t}, \mathbf{x}) \tag{2.8}$$

Where y_t is the token at step t and $\mathbf{y}_{<t}$ represents all preceding tokens $y_1, \ldots, y_{t-1}$ in the generated sequence. To compute these probabilities, the Transformer processes the input $\mathbf{x}$ and the generated history $\mathbf{y}_{<t}$ to produce a vector of non-normalized scores, known as logits, for every token in the vocabulary $\mathcal{V}$. Let z_i be the logit corresponding to the i-th token $v_i \in \mathcal{V}$. The model converts these logits into a valid probability distribution using the Softmax function [Bri90; Vas+17]:

$$P(y_t = v_i \mid \mathbf{y}_{<t}, \mathbf{x}) = \frac{\exp(z_i)}{\sum_{j=1}^{|\mathcal{V}|} \exp(z_j)} \qquad \forall i \in 1, \ldots, |\mathcal{V}| \tag{2.9}$$

This process is repeated iteratively, where the token selected at step t is appended to the input for step $t+1$ [Vas+17].

2.3.2 Decoding Strategies and the Role of Temperature

Once the probability distribution over the Vocabulary is computed for the next token, a **decoding strategy** determines what token will be selected. Common maximization-based strategies include **Greedy Decoding** or **Beam Search**. **Greedy Decoding** is the simplest deterministic method. It selects the token with the highest probability mass at the current time step t [SVL14]:

$$y_t = \arg\max_{y \in \mathcal{V}} P(y \mid \mathbf{y}_{<t}, \mathbf{x}) \tag{2.10}$$

While computationally efficient, this strategy often yields sub-optimal results. By focusing solely on local probability maximization, the strategy risks settling into local optima, effectively missing high-probability sequences hidden behind lower-probability initial tokens. Furthermore, greedy decoding is known to produce repetitive and generic text, as it fails to account for the diversity and global coherence of the sequence [WR16; Hol+20].

Beam Search extends Greedy Decoding by maintaining multiple hypotheses at each time step, rather than a single sequence. It introduces a hyperparameter k, known as the **beam width**, which determines the number of candidates to keep. At each step t, the algorithm expands all k candidates from the previous step by appending every possible next token from the vocabulary $\mathcal{V}$ [FA17].

Formally, let $\mathcal{B}_{t-1} = \{\mathbf{y}_{<t}^{(1)}, \ldots, \mathbf{y}_{<t}^{(k)}\}$ be the set of k hypothesis sequences generated up to step $t-1$. The algorithm proceeds in three steps [FA17]:

1. **Expansion**: Each hypothesis in $\mathcal{B}_{t-1}$ is extended with every possible token v from the vocabulary $\mathcal{V}$, creating a candidate set $\mathcal{C}_t$.

2. **Scoring**: Each candidate sequence is assigned a score based on its cumulative log-probability. The score for a sequence $\mathbf{y}_{1:t}$ (representing $y_1, \ldots, y_t$) is defined recursively:

$$\text{Score}(\mathbf{y}_{1:t}) = \text{Score}(\mathbf{y}_{1:t-1}) + \log P(y_t \mid \mathbf{y}_{<t}, \mathbf{x}) \tag{2.11}$$

3. **Selection (Pruning)**: The algorithm selects the k candidates with the highest scores to form the new hypothesis set $\mathcal{B}_t$:

$$\mathcal{B}_t = \underset{\mathbf{y}_{1:t} \in \mathcal{C}_t}{\text{arg-top-}k} \left(\text{Score}(\mathbf{y}_{1:t})\right) \tag{2.12}$$

This approach allows the model to recover from locally sub-optimal decisions if they lead to a highly probable global sequence. However, Beam Search is computationally more expensive than greedy methods and, as noted by Holtzman et al. [Hol+20], can still suffer from repetitive loops or generic responses in open-ended generation tasks.

Sampling represents a fundamental shift from maximization-based decoding strategies towards stochastic generation. Instead of strictly selecting the token with the highest probability, the model randomly draws the next token y_t from the conditional probability distribution $P(y_t \mid \mathbf{y}_{<t}, \mathbf{x})$.

Formally, at each time step t, the next token is sampled according to:

$$y_t \sim P(y \mid \mathbf{y}_{<t}, \mathbf{x}) \tag{2.13}$$

This stochasticity allows the model to generate more diverse and creative outputs for the same input. However, pure random sampling from the full distribution can occasionally lead to incoherence. If the "tail" of the distribution, that contains very low-probability tokens, is sampled, it may disrupt the semantic flow or factual accuracy of the sequence. To mitigate this risk while preserving diversity, two truncation techniques are commonly employed:

1. **Top-k Sampling** [FLD18]: The model samples from the top k most probable tokens. This effectively truncates the unreliable tail of the distribution.

2. **Nucleus (Top-p) Sampling** [Hol+20]: The model samples from the smallest set of tokens whose cumulative probability exceeds a threshold p (e.g., $p = 0.9$). Unlike Top-k, which uses a fixed number of tokens, Top-p adapts dynamically. The set of candidates grows when the model is uncertain (flat distribution) and shrinks when the model is confident (peaked distribution).

While these truncation methods improve coherence, they operate by removing options. A more continuous method for controlling the shape of the probability distribution, and the core mechanism for this thesis, is **Temperature Scaling**.

Temperature Scaling is a technique used to calibrate the confidence of the model's predictions. Unlike Top-k or Nucleus Sampling, which truncate the vocabulary, temperature scaling modifies the probability distribution itself by rescaling the logits before the Softmax function is applied [HVD+15].

Let z_i be the logit for token v_i. We introduce a hyperparameter $\tau > 0$, called **temperature**. The probability of selecting token v_i is then given by:

$$P^{\tau}(y_t = v_i \mid \mathbf{y}_{<t}, \mathbf{x}) = \frac{\exp(z_i/\tau)}{\sum_{j=1}^{|\mathcal{V}|} \exp(z_j/\tau)} \tag{2.14}$$

The value of τ controls the entropy of the output distribution. When choosing a **low temperature**, $\tau < 1.0$, the distribution becomes sharper, i.e., more peaked. The difference between large and small logits is exaggerated, causing the model to become more confident and deterministic. In the limit, this approximates greedy decoding. Using a **high temperature**, $\tau > 1.0$, flattens the distribution towards a uniform distribution. This increases the likelihood of selecting lower-probability tokens, resulting in more diverse but also potentially more incoherent or hallucinatory outputs. For the temperature $\tau = 1.0$, the distribution remains unchanged, reflecting the model's original logits [HVD+15].

In standard generation tasks, τ is typically a fixed hyperparameter tuned to balance diversity and equality. However, in this thesis, we interpret temperature variation as a mechanism to explore the model's epistemic uncertainty.

2.4 Quantifying Uncertainty in LLMs

Applying the principles of uncertainty quantification to LLMs presents unique challenges that do not exist in standard classification tasks. In classification, the label space is fixed and mutually exclusive. In natural language generation, however, the model must navigate a massive, interconnected vocabulary over multiple autoregressive steps, making it difficult to disentangle aleatoric from epistemic uncertainty.

In NLP, aleatoric uncertainty is highly prevalent due to the inherent ambiguity of language. Often, there are multiple semantically equivalent ways to express the same underlying thought. For example, "It is raining heavily" and "It is pouring outside" describe the same state, distributing the model's probability mass across different valid token sequences [KGF23].

Epistemic uncertainty arises when the model suffers from a genuine lack of knowledge, such as when it faces out-of-distribution (OOD) prompts or queries about facts absent from its training data [KG17; HW21].

2.4.1 Negative Log-Likelihood in LLMs

While standard classification evaluates the likelihood of a single class prediction, evaluating an LLM requires quantifying the model's confidence over an entire generated sequence. For a given input context $\mathbf{x}$ and a target sequence $\mathbf{y}$, the sequence likelihood is the product of the individual token probabilities, as defined by the chain rule (Equation 2.8).

Because multiplying many small probabilities for long sequences leads to numerical underflow, it is standard practice to compute the **Negative Log-Likelihood (NLL)**

[GBC16]. Taking the negative logarithm transforms the product of the probabilities into a sum. Furthermore, to fairly compare sequences of different lengths, the NLL can be normalized by the total number of tokens T, yielding the average NLL per token [MC18]:

$$\text{NLL}(\mathbf{y} \mid \mathbf{x}) = -\sum_{t=1}^{T} \log P(y_t \mid \mathbf{y}_{<t}, \mathbf{x}) \tag{2.15}$$

$$\text{average NLL}(\mathbf{y} \mid \mathbf{x}) = -\frac{1}{T}\sum_{t=1}^{T} \log P(y_t \mid \mathbf{y}_{<t}, \mathbf{x}) \tag{2.16}$$

A lower NLL indicates that the model assigns a higher probability to the target sequence, implying higher confidence.

2.5 Relative Likelihood and Valid Temperature Sets

To systematically construct a credal set, we need a rigorous criterion to distinguish between plausible and implausible probability distributions. Following recent advances in credal prediction [Löh+25], we achieve this by employing the statistical concept of **relative likelihood** [Wil38; CH79].

2.5.1 Relative Likelihood

In standard maximum likelihood estimation, the goal is to find the single best parameter that maximizes the likelihood of the observed data. In the context of our method, we treat the decoding temperature $\tau \in \mathcal{T}$ as the parameter of interest. Let $L(\tau)$ denote the likelihood of the target sequence generated at temperature τ. Furthermore, let τ^{ML} be the optimal temperature that yields the highest likelihood (or equivalently, minimizes the NLL). To characterize the plausibility of any other temperature τ, we compute its relative likelihood $\gamma(\tau)$, defined as the ratio of its likelihood to the maximum likelihood [Wil38; Bir62; Was90; WM99]:

$$\gamma(\tau) = \frac{L(\tau)}{L(\tau^{ML})} = \frac{L(\tau)}{\sup_{\tau' \in \mathcal{T}} L(\tau')} \tag{2.17}$$

By definition, $\gamma(\tau)$ is bounded in the interval $[0, 1]$, where $\gamma(\tau^{ML}) = 1$. The relative likelihood provides an intuitive measure of how well a specific temperature explains the data compared to the optimal temperature setting.

2.5.2 Constructing α-Cuts for Temperatures

On the basis of the relative likelihood, a set of valid temperatures can be constructed by including only those that are plausible in the sense of surpassing a specific threshold $\alpha \in [0, 1]$. This set is formally referred to as an α**-cut** [ACC12]:

$$\mathcal{T}_\alpha = \{\tau \in \mathcal{T} : \gamma(\tau) \geq \alpha\} \tag{2.18}$$

According to this definition, a temperature τ is considered implausible, and is therefore excluded from the valid set, if its likelihood is too small compared to the likelihood of the best temperature, specifically if it is less than α times the optimal likelihood.

Related Work

3

The reliable quantification of uncertainty in LLMs has garnered significant attention in recent years. As researchers seek to mitigate hallucinations and improve model safety, various metrics have been proposed to estimate when a model's generation is likely to be incorrect. This chapter reviews the evolution of these methods, starting from traditional UQ metrics, and culminating in state-of-the-art approaches like Semantic Entropy, which serves as the primary baseline for the credal prediction methodology developed in this thesis.

Applying traditional UQ techniques to LLMs presents substantial challenges. The computational overhead of Deep Ensembles [LPB17] or Bayesian frameworks like Monte Carlo dropout [GG16] makes them impractical for natural language generation (NLG) in massive models. Initial approaches to estimating LLM uncertainty directly adapted standard classification metrics or relied on heuristic comparisons. These baselines generally fall into three categories: probability-based methods, lexical similarity and self-evaluation.

3.1 Probability-Based Methods

The most straightforward methods evaluate the model's confidence using the probability distributions generated during decoding. Token-level entropy evaluates local uncertainty but fails to capture the holistic confidence of a factual claim. To address this, Malinin and Gales [MG21] proposed sequence-level predictive entropy (often implemented as length-normalized entropy), which estimates uncertainty across all possible output sequences.

Another prominent approach within this category is the **Margin Measure** [LLS22]. Instead of calculating the entropy over the entire distribution, the margin measure isolates the gap in log-probability between the model's most likely sequence and its second most likely sequence. A small margin implies that the model is highly conflicted between its top two choices, indicating high uncertainty.

However, both predictive entropy and the margin measure, operate strictly on exact token matches, they artificially inflate uncertainty when an LLM distributes probability mass across multiple valid phrasings of the same answer.

3.2 Lexical Similarity

To account for the fact that generated answers might differ slightly in their surface forms, researchers adapted translation metrics to quantify uncertainty. By sampling multiple answers for a given question, one can compute the average lexical similarity between them using metrics like ROUGE-L [Lin04] or BLEU [Pap+02]. If the sampled answers have high n-gram overlap, the model is deemed certain. While an improvement over strict sequence likelihood, lexical similarity still struggles with open-ended generation. An LLM might generate "Berlin" and "The capital of Germany", which share zero lexical overlap (low ROUGE-L) but convey the exact same meaning, leading to false uncertainty signals.

3.3 Self-Evaluation

Another distinct approach treats the LLM as an independent evaluator of its own uncertainty. Methods like P(True) [Kad+22] prompt the model to evaluate its own generated answer or ask it directly if a proposed fact is correct. While intuitive, self-evaluation baselines are highly sensitive to the specific phrasing of the prompt.

All of these baseline families fail to reliably distinguish between benign linguistic flexibility and genuine factual ignorance. Traditional metrics treat these lexical variations as disjoint, competing hypotheses. This critical limitation necessitated a paradigm shift from lexical (word-based) probability to semantic (meaning-based) probability, leading to the development of Semantic Entropy [KGF23].

3.4 Semantic Entropy

To resolve the conflation of lexical flexibility and factual ignorance, Kuhn et al. [KGF23] introduced the concept of **Semantic Entropy (SE)**. Semantic entropy shifts the unit of uncertainty measurement from the exact sequence of words to the underlying meaning, or *semantics*, of the generated text. Because SE serves as both the primary baseline and a core mechanical component of the methodology proposed in this thesis, we detail its formal computation below. The SE itself is computed in three steps:

1. Generating N sequences by the same model
2. Clustering the generated sequences by semantic equivalence
3. Computation of the semantic entropy

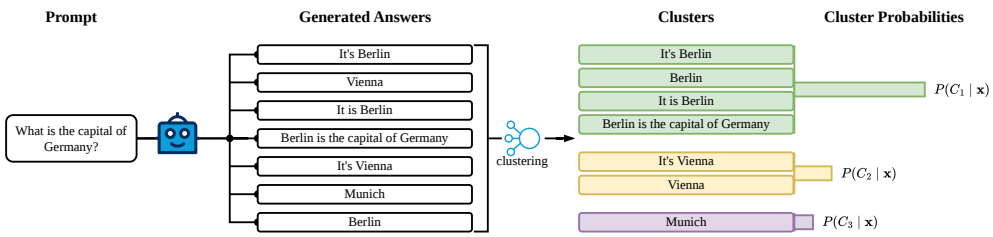

Fig. 3.1.: Visualization of semantic clustering for the question "What is the capital of Germany?". After the generation of seven answer strings, they are semantically clustered and afterwards the the cluster probabilities are computed as shown in 3.1.

3.4.1 Semantic Clustering

The core premise of SE is that the infinite space of possible lexical sequences generated by an LLM can be partitioned into a finite set of discrete meanings.

Let $\mathbf{x}$ be the input prompt, and let $\mathcal{S} = \{s^{(1)}, s^{(2)}, \ldots, s^{(N)}\}$ be a set of N sequences generated by the model. Kuhn et al. [KGF23] define a bidirectional semantic equivalence relation between sequences, denoted as $E(\cdot, \cdot)$, which holds true if sequence $s^{(i)}$ and sequence $s^{(j)}$ convey the exact same factual information, i.e., they entail, in the context of the prompt $\mathbf{x}$. The sequences in a set $c \in \mathcal{C}$, where $\mathcal{C}$ is the space of all possible semantic equivalence classes, all share the same meaning, i.e., $\forall s, s' \in c : E(s, s')$.

This equivalence relation groups the generated sequences into distinct **semantic clusters** (or equivalence classes), denoted as $C = \{c_1, c_2, \ldots, c_K\}$, where $\forall s^{(i)}, s^{(j)} \in c_k : E(s^{(i)}, s^{(j)})$ and $K \leq N$. In practice, this clustering is performed using an automated evaluation model, such as a Natural Language Inference (NLI) classifier (Kuhn et al. [KGF23] used the DeBERTa-large-MNLI model [He+21]).

3.4.2 Computing the Semantic Entropy

Once the sequences are clustered, the probability of a specific meaning (a cluster $c \in \mathcal{C}$) is computed by marginalizing over the probabilities of all individual lexical sequences within that cluster. The cluster probability is given by:

$$P(c \mid \mathbf{x}) = \sum_{s \in c} P(s \mid \mathbf{x}) \tag{3.1}$$

By treating the clusters as the new categorical variables, the SE is calculated as the Shannon Entropy over the cluster probabilities:

$$\mathrm{SE}(\mathbf{x}) = -\sum_{c} P(c \mid \mathbf{x}) \log P(c \mid \mathbf{x}) \tag{3.2}$$

Because we can not know every possible meaning-class c, we can only sample c from the given distribution by the model. Thus Kuhn et al. [KGF23] use Monte Carlo integration over the semantic equivalence classes to estimate the expectation in Equation 3.2:

$$\mathrm{SE}(\mathbf{x}) \approx -|C|^{-1} \sum_{i=1}^{|C|} \log P(C_i \mid \mathbf{x}) \tag{3.3}$$

where

$$P(C_i \mid \mathbf{x}) = \frac{P(c_i \mid \mathbf{x})}{\sum_{c \in C} P(c \mid \mathbf{x})} \tag{3.4}$$

SE represents a significant advancement over raw sequence likelihood. But it relies on a single, static probability distribution. In the standard implementation by Kuhn et al. [KGF23], sequences are sampled, and probabilities are evaluated at a single, fixed temperature (Kuhn et al. [KGF23] state that $\tau = 0.5$ is optimal).

As established in Section 2.2, precise probabilities are brittle when a model is miscalibrated or highly ignorant. This gap motivates the primary contribution of this thesis. Instead of computing SE on a single distribution, we propose computing it across a **credal set of distributions** generated by a valid range of temperatures. By analyzing how semantic cluster probabilities fluctuate across the temperatures, yielding upper and lower cluster bounds, we detect epistemic uncertainty that standard SE is blind to.

3.5 Credal Prediction Based on Relative Likelihood

While SE addresses the issue of lexical equivalence in LLMs, it lacks a mechanism to handle severe epistemic uncertainty arising from a poorly calibrated, single probability distribution. To address this, we look to recent advancements in the field of imprecise probabilities, specifically the application of credal sets to deep learning.

3.5.1 Relative Likelihood Ensembles

Recently, Löhr et al. [Löh+25] proposed a theoretically grounded approach to credal prediction based on the statistical notion of relative likelihood. Instead of relying on a single maximum likelihood estimator (MLE), their framework defines the prediction target as a credal set induced by all plausible models whose relative likelihood exceeds a specific threshold α (the α-cut). Specifically, they introduce the Credal Relative Likelihood ($CreRL_\alpha$) method, which trains an ensemble of neural networks using an early-stopping strategy. Each model in the ensemble is trained until it reaches a specific relative likelihood threshold τ_i, ensuring that the final ensemble spans the entire valid α-cut. To guarantee diversity among the ensemble members, they further employ a novel "ToBias" initialization strategy, which initializes the models to predict degenerate distributions at the vertices of the probability simplex.

3.5.2 From Ensembles to LLM Temperatures

The $CreRL_\alpha$ framework successfully demonstrates that relative likelihood can be used to construct credal sets, achieving superior coverage on standard classification benchmarks. However, this approach relies on optimizing model weights during training. Applying this weight-space ensemble approach directly to modern LLMs is computationally prohibitive. Training multiple copies of a multi-billion parameter model to convergence is infeasible for most applications. It also does not cleanly translate to the autoregressive, open-ended generation of sequences.

This limitation motivates the core innovation of our thesis: shifting the UQ from the *weight space* (ensembling models) to the *decoding space* (ensembling temperature).

4 Methodology: Credal Semantic Entropy via Temperature Scaling

This chapter introduces a novel methodology that constructs a credal set for LLMs by dynamically scaling the decoding temperature. By treating the temperature hyperparameter as a continuous parameter space and applying the statistical principles of relative likelihood, we systematically filter out implausible distributions. We then extend semantic entropy to operate over this credal set, yielding upper and lower uncertainty bounds that serve as a signal for high model uncertainty.

4.1 Overview of the Pipeline

Our proposed methodology, which we term Credal Semantic Entropy, works for a given input prompt $\mathbf{x}$ (a question and optional context), in four sequential stages:

1. **Valid Temperature Selection (The α-cut)**: Before generation, we must determine which temperatures produce valid probability distributions in the context of relative likelihood. Using a labeled dataset, we evaluate the model's average Negative Log-Likelihood (NLL) as shown in Equation 2.16 across a discrete grid of temperatures. We identify the optimal temperature that minimizes the average NLL and compute the relative likelihood for all temperatures. The α-cut is constructed by all temperatures with a relative likelihood greater or equal than α.

2. **Generation and Semantic Clustering**: Using the prompt $\mathbf{x}$, the LLM generates multiple candidate responses at every valid temperature in the α-cut. We cluster these generated sequences by their meaning over all valid temperatures.

3. **Computation of Cluster Probability Bounds**: For each valid temperature, we compute the probability of each semantic cluster. Because the temperature modulates the entropy of the underlying token distributions, the probability of a given cluster will fluctuate across the α-cut. By finding the minimum and maximum probability assigned to each cluster across all valid temperatures, we establish the lower and upper bounds for every semantic concept.

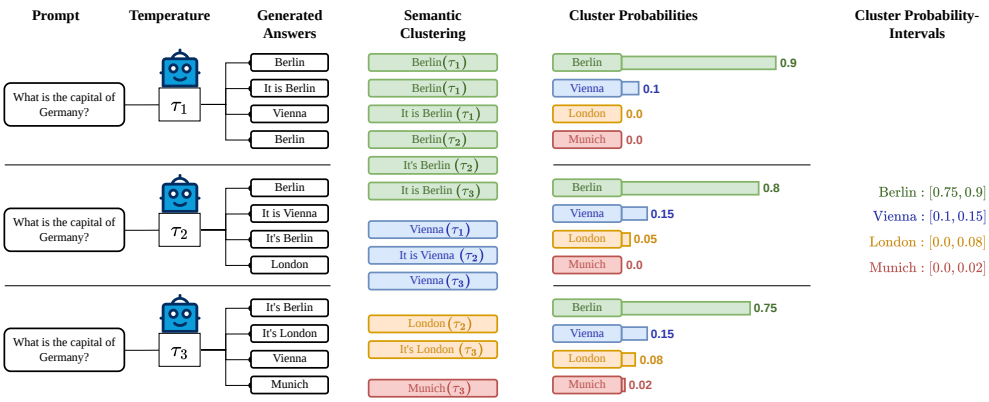

Fig. 4.1.: Visualization of the first three stages of our methodology. The process begins with multi-temperature generation, followed by semantic clustering performed on the union of all generated sequences. Finally, the probability mass of each cluster is marginalized at every temperature to establish the probability intervals used for UQ.

4. **Quantifying Uncertainty**: Finally, we compute the Upper Shannon Entropy derived from these probability bounds. By maximizing the entropy within the credal set, we capture the overall model uncertainty. In the context of open-domain question-answering, this resulting scalar serves as a conservative confidence metric to predict the factual reliability of the generated answer.

The following sections formally define the mathematical and computational mechanics of each stage in this pipeline.

4.2 Valid Temperature Selection (the $\alpha-$Cut)

We consider the task of free-form Question Answering (QA) using a pre-trained, autoregressive LLM $\mathcal{M}$ with a vocabulary $\mathcal{V}$. This process is performed a priori using a labeled training dataset or calibration dataset. We evaluate a predefined grid of candidate temperatures $\mathcal{T}$. For each temperature $\tau \in \mathcal{T}$, we measure its ability to assign high probability to the ground-truth answers. This involves three steps: computing the average NLL, deriving the relative likelihood and applying an α-cut threshold.

4.2.1 Computing the Average Negative Log-Likelihood

For every question-answer pair $(\mathbf{x}^{(i)}, \mathbf{y}^{(i)})$ in the training dataset $\mathcal{D}_{\text{Train}}$ of size D, we compute the length-normalized NLL (2.16) of the ground-truth answer sequence

under the temperature-scaled probability distribution P^τ. The average NLL for a given temperature τ across the entire dataset is computed as:

$$\overline{\text{NLL}}(\tau) = \frac{1}{D}\sum_{i=1}^{D}\left(-\frac{1}{T_i}\sum_{t=1}^{T_i}\log P^\tau(y_t^{(i)} \mid \mathbf{y}_{<t}^{(i)}, \mathbf{x}^{(i)})\right) \tag{4.1}$$

where T_i is the length (number of tokens) of the i-th target sequence. This average NLL serves as our primary metric for evaluating how well a specific temperature calibrates the model's output distribution to the true data distribution. We denote the optimal temperature that yields the minimum $\overline{\text{NLL}}$ as τ^{ML}:

$$\tau^{ML} = \arg\min_{\tau\in\mathcal{T}} \overline{\text{NLL}}(\tau) \tag{4.2}$$

4.2.2 Deriving Relative Likelihoods

Because NLL values are absolute and can vary drastically depending on the model and dataset, it is difficult to define a universal threshold for what constitutes a "good" temperature. To standardize this, we map the average NLLs into the relative likelihood space.

As established in Equation 2.17, the relative likelihood $\gamma(\tau)$ of any temperature is computed by comparing its likelihood $L(\tau)$ to the optimal likelihood $L(\tau^{ML})$. In the logarithmic domain, this is equivalent to the exponentiated difference of their average NLLs:

$$\gamma(\tau) = \frac{L(\tau)}{L(\tau^{ML})} = \exp\left(-\left(\overline{\text{NLL}}(\tau) - \overline{\text{NLL}}(\tau^{ML})\right)\right) \tag{4.3}$$

By definition, the optimal temperature τ^{ML} has a relative likelihood of $\gamma(\tau^{ML}) = 1.0$. All other temperatures τ will have a relative likelihood $\gamma(\tau) \in [0, 1]$, explicitly quantifying how plausible they are compared to the best performing parameterization.

4.2.3 The α-Cut

Following the framework of credal prediction [Löh+25], we apply an α-cut (Section 2.5.2), retaining only those temperatures whose relative likelihood meets or exceeds a specified confidence threshold $\alpha \in [0, 1]$. The final set of valid temperatures is formally defined as:

$$\mathcal{T}_\alpha = \{\tau \in \mathcal{T} : \gamma(\tau) \geq \alpha\} \tag{4.4}$$

4.3 Generation and Semantic Clustering

Once the set of valid temperatures $\mathcal{T}_\alpha$ has been established via the training dataset, we apply the methodology to the unseen test queries. For a given input prompt $\mathbf{x}$, the goal is to map the model's raw token outputs into discrete semantic concepts. This requires a two-step process. Generating a diverse pool of candidate responses and clustering them based on semantic equivalence.

4.3.1 Multi-Temperature Sequence Generation

To evaluate the model's uncertainty, we must sample responses from every valid temperature. For each temperature $\tau \in \mathcal{T}_\alpha$, we prompt the LLM $\mathcal{M}$ to generate N independent answer sequences.

Let $\mathcal{S}_\tau = \{s_1^\tau, s_2^\tau, \ldots, s_N^\tau\}$ denote the set of sequences generated at temperature τ. The total pool of generated responses for the prompt $\mathbf{x}$ is the union of all sequences generated across all valid temperatures:

$$\mathcal{S}_{total} = \bigcup_{\tau \in \mathcal{T}_\alpha} \mathcal{S}_\tau \tag{4.5}$$

By generating across $\mathcal{T}_\alpha$, we force the model to explore different regions of its probability space. Lower temperatures favor the most likely tokens, whereas higher temperatures flatten the distribution, encouraging the model to output lower-probability claims.

4.3.2 Bidirectional Entailment and Equivalence Classes

As discussed in Section 3.4, evaluating uncertainty on the raw lexical set $\mathcal{S}_{total}$ artificially inflates misconfidence due to the aleatoric linguistic variations. We therefore partition $\mathcal{S}_{total}$ into a set of semantic equivalence classes, or clusters, denoted as $C = \{c_1, c_2, \ldots, c_K\}$ by pairwise comparing all sequences across all valid temperatures as described in the following. To determine if two distinct sequences s_i and s_j belong to the same cluster, we employ an automated Natural Language Inference (NLI) model (e.g., DeBERTa-large-MNLI [He+21]). The NLI

model evaluates whether the premise sequence entails the hypothesis sequence. We define a strict semantic equivalence relation $E(\cdot,\cdot)$ which holds true if and only if there is **bidirectional entailment**: $(s_i \models s_j) \wedge (s_j \models s_i)$

If s_i entails s_j and s_j entails s_i, the sequences are deemed factually equivalent in the context of the prompt $\mathbf{x}$ and are assigned to the same cluster c_k. Thus, a cluster holds all sequences that belong to the same equivalence class of $E(\cdot,\cdot)$: $\forall s, s' \in c_k : E(s, s')$.

4.3.3 Constructing the Cluster Probability Distributions

Once the clusters C are established, we must compute the probability of each cluster under each temperature $\tau \in \mathcal{T}_\alpha$. The probability of a semantic cluster c_k at temperature τ is the normalized sum of the likelihoods of all sequences within that cluster, evaluated under the distribution P^τ:

$$P^\tau(C_k \mid \mathbf{x}) = \frac{P^\tau(c_k \mid \mathbf{x})}{\sum_{c \in C} P^\tau(c \mid \mathbf{x})} \tag{4.6}$$

where

$$P^\tau(c_k \mid \mathbf{x}) = \sum_{s \in c_k} P^\tau(s \mid \mathbf{x}) \tag{4.7}$$

This marginalization step is crucial. It abstracts away the specific phrasing of the answers, yielding a set of discrete probability distributions $P^\tau(C)$ over the fundamental semantic meaning. It is important to note that the set of semantic clusters C is global. It contains every unique meaning generated across *all* valid temperatures. If a specific cluster c_k does not contain any sequences generated at a particular temperature τ, we assign it a zero probability mass for that distribution. This ensures that every distribution P^τ is defined over the same common support C, allowing for the direct element-wise comparison required to compute the upper and lower probability bounds in the next step. To compute the sequence probability $P^\tau(s \mid \mathbf{x})$, we use the length normalized NLL as defined in Equation 2.16 because we do not want to punish longer sequences [MG21]. In the case where shorter answers are wanted, as it is in TriviaQA [Jos+17], we could use the unnormalized NLLs (2.15) as well.

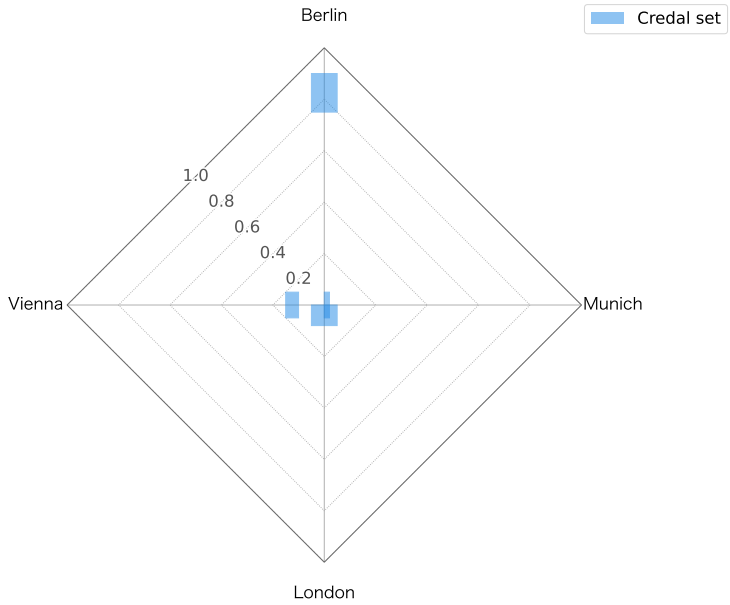

Fig. 4.2.: Geometric representation of the credal set derived from the probability intervals established in Figure 4.1. Each axis of the radar chart corresponds to a distinct semantic cluster. The blue bars represent the probability intervals (lower to upper bounds) for each class, illustrating the model's epistemic uncertainty across the potential predictions. This visualization approach follows [Hof+26].

4.3.4 Quantifying Credal Semantic Uncertainty

By scaling the temperature, we observe a range of probabilities for each semantic cluster. We utilize the observed probability fluctuations to establish upper and lower probability bounds for every semantic concept. The **credal set** is then formally defined as the set of *all* valid probability distributions that respect these element-wise bounds.

Establishing Probability Bounds
Let $C = \{c_1, c_2, \ldots, c_K\}$ be the set of identified semantic clusters. For each cluster c_k, we inspect its probability mass across all valid temperatures $\tau \in \mathcal{T}_\alpha$ We define the **lower probability** $\underline{P}(c_k \mid \mathbf{x})$ and the **upper probability** $\overline{P}(c_k \mid \mathbf{x})$ as the minimum and maximum probabilities observed for that cluster.

$$\underline{P}(c_k \mid \mathbf{x}) = \min_{\tau \in \mathcal{T}_\alpha} P^\tau(C_k \mid \mathbf{x}) \qquad \overline{P}(c_k \mid \mathbf{x}) = \max_{\tau \in \mathcal{T}_\alpha} P^\tau(C_k \mid \mathbf{x}) \tag{4.8}$$

These bounds define an interval $[\underline{P}(c_k \mid \mathbf{x}), \overline{P}(c_k \mid \mathbf{x})]$ representing the model's epistemic ambiguity regarding the plausibility of meaning c_k.

Defining the Credal Set
We construct the credal set $\mathcal{K}$ as the convex set of all probability distributions P over C that are consistent with these interval bounds. Formally:

$$\mathcal{K} = \left\{ P \in \Delta_K \mid \forall k : \underline{P}(c_k \mid \mathbf{x}) \leq P(c_k \mid \mathbf{x}) \leq \overline{P}(c_k \mid \mathbf{x}), \sum_{k=1}^{K} P(c_k \mid \mathbf{x}) = 1 \right\} \tag{4.9}$$

where Δ_K denotes the standard probability simplex. By defining $\mathcal{K}$ via these bounds, we capture not just the specific distributions observed at sampled temperatures, but the entire region of probability space supported by the valid relative likelihoods.

Uncertainty Quantification
To measure the uncertainty over the semantic clusters, we adapt the credal evaluation framework from standard classification tasks, as detailed in Section 2.2.3. By treating each semantic cluster as a distinct class, we can compute the lower and upper bounds of the Shannon Entropy to quantify the model's uncertainty.

Upper Semantic Entropy (H^*)
This metric corresponds to the maximum entropy (highest uncertainty) attainable by any distribution within the credal set. It serves as our primary signal for high model uncertainty (2.2.3), identifying cases where the model's interval constraints allow for a highly flat, uniform distribution:

$$H^*(\mathcal{K}) = \max_{P \in \mathcal{K}} \left(-\sum_{k=1}^{K} P(c_k \mid \mathbf{x}) \log P(c_k \mid \mathbf{x}) \right) \tag{4.10}$$

Lower Semantic Entropy (H_*)
Conversely, the lower entropy corresponds to the most deterministic distribution consistent with the bounds:

$$H_*(\mathcal{K}) = \min_{P \in \mathcal{K}} \left(-\sum_{k=1}^{K} P(c_k \mid \mathbf{x}) \log P(c_k \mid \mathbf{x}) \right) \tag{4.11}$$

The difference $H^* - H_*$ reflects the magnitude of epistemic uncertainty as shown in Section 2.2.3. For the experimental evaluation in this thesis, we utilize the upper semantic entropy H^*, as it provides the most conservative safety guarantee by considering the maximum plausible uncertainty. This makes it ideally suited for evaluating a model's answers in a conversational context. In such settings, our primary goal is to bound the total predictive uncertainty to ensure safe responses, rather than explicitly distinguish between the model's lack of knowledge and inherent data noise.

5 Experiments

To validate the effectiveness of our proposed Credal Semantic Entropy method, we conducted a series of experiments on a standard free-form Question Answering (QA) benchmark. The primary objective of these experiments is to evaluate whether quantifying uncertainty via valid temperatures provides a more reliable signal for high uncertainty than standard single-point metrics.

5.1 Experimental Setup

5.1.1 Dataset: CoQA

We utilized the **Conversational Question-Answering (CoQA)** dataset [RCM19] for our evaluation. CoQA is a large-scale dataset for building Conversational Question Answering systems. It contains 127k+ questions with answers obtained from 8k+ conversations about text passages from seven diverse domains [RCM19].

For our specific task, we followed the experimental protocol established by Kuhn et al. [KGF23]. We used the development split of CoQA, which provides a diverse set of questions. The dataset was split into a **calibration set** (used to determine the valid temperature set) and a **test set** (used to evaluate the final uncertainty metrics).

5.1.2 Model Architecture

All experiments were conducted using the **Opt-6.7B** model [Zha+22]. OPT (Open Pre-trained Transformer) is a decoder-only architecture that replicates the performance and behaviour of GPT-3 [Bro+20]. We chose the 6.7B parameter variant as it strikes a balance between computational feasibility and the capability to generate coherent text.

For the semantic clustering step, we employed a **DeBERTa-large-MNLI** model [He+21] to assess bidirectional entailment between generated answers.

5.1.3 Baselines

We compared our proposed **Credal Semantic Entropy (Credal SE)** against several strong baselines:

- **Semantic Entropy (SE)**: The original metric proposed by Kuhn et al. [KGF23],
- **Lexical Similarity (ROUGE-L)**: We compute the average ROUGE-L score across all pairs of the N sampled sequences. Lower average similarity indicates higher lexical dispersion (uncertainty). [Lin04].
- **Predictive Entropy (PE)**: The standard entropy of the model's output distribution over the sequence of tokens. Unlike semantic entropy, this metric considers every lexical variation as a distinct outcome, calculating uncertainty directly on the raw strings without semantic clustering [MG21].
- **Length-Normalized PE**: A variant of Predictive Entropy that normalizes the sum of token entropies by the sequence length, mitigating the bias against longer answers [MG21].
- **Margin Measure (MM)**: The difference in log-probability between the most likely sequence and the second most likely sequence. A small margin indicates high uncertainty between the top two choices [LLS22].
- **P(true)**: A self-evaluation metric [Kad+22] where the model is prompted with the question and its generated answer, and asked to output "True" or "False". We extract the probability mass assigned to the "True" token.

5.1.4 Evaluation Metric: AUROC

To quantify the performance of each uncertainty method, we treat uncertainty quantification (UQ) as a binary classification problem. The core task is to determine if the uncertainty score can reliably distinguish between correct and incorrect model generations. We use the answers that were generated by the model using the temperature with the maximum likelihood.

Ground Truth Labeling
To establish the ground-truth labels for our experiments, we compare the generated answers against the gold-standard reference answer provided in the CoQA dataset. We define correctness based on lexical overlap using the **ROUGE** metric (Recall-Oriented Understudy for Gisting Evaluation) [Lin04]. Specifically, a generated answer is labeled as:

- **Correct**: If the ROUGE-L score between the generation and the reference is strictly greater than 0.3.

- **Incorrect**: If the ROUGE-L score is less than or equal to 0.3.

This threshold allows for minor phrasing differences while ensuring that the core semantic content overlaps significantly with the reference.

AUROC Calculation
We report the **Area Under the Receiver Operating Characteristic curve (AUROC)** [Bra97]. The AUROC represents the probability that a randomly selected incorrect answer will have a higher uncertainty score than a randomly selected correct answer.

- AUROC $= 0.5$: Indicates random performance (no better than a coin flip)

- AUROC $= 1.0$: Indicates perfect detection (all incorrect answers have higher uncertainty scores than all correct answers).

5.2 Results

Constructing a credal set via multi-temperature-sampling incurs significant computational overhead. To maintain a feasible evaluation pipeline while ensuring sufficient granularity, we restricted our analysis to the temperature interval $[0.4, 2.2]$ evaluated at two distinct resolutions: a 20-step grid (Table A.1) and a 50-step grid (Table A.2).

Our evaluation is conducted on a subset of approximately 2400 queries, using 60% (≈ 1440 questions) of the questions to find valid temperatures and the other 40% (≈ 960 questions) to evaluate our uncertainty estimation, from the CoQA development set. In this section, we systematically present the UQ performance (measured using AUROC) across these configurations, exploring how the grid resolution and the relative likelihood threshold α impact the robustness of our uncertainty estimates.

5.2.1 Evaluation: 20-Step Interval

We begin our analysis using the coarser 20-step temperature interval combined with a highly restrictive relative likelihood threshold of $\alpha = 0.9$. Table 5.1 summarizes the AUROC scores for our proposed Credal Semantic Entropy alongside the established single-point baselines. The baselines were evaluated on the generations predicted by the model using the most likely temperature ($\tau = 1.347$).

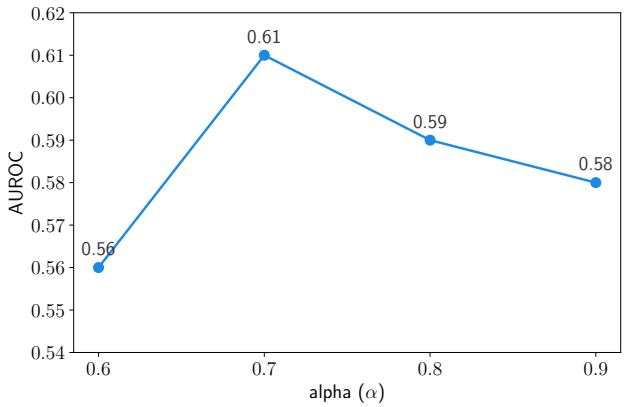

Fig. 5.1.: Impact of the relative likelihood threshold (α) on the UQ performance of Credal Semantic Entropy. The line plot illustrates the AUROC scores evaluated on the 20-step temperature grid for $\alpha \in \{0.6, 0.7, 0.8, 0.9\}$.

Table 5.1: UQ performance (AUROC) on the CoQA subset using a 20-step temperature interval and a strict relative likelihood threshold ($\alpha = 0.9$). Metrics compared are our proposed Credal Semantic Entropy (Credal SE), standard Semantic Entropy (SE), Predictive Entropy (PE), length-normalized Predictive Entropy (ln PE), Margin Measure (MM), P(true) and Lexical Similarity (LS).

Configuration	Credal SE	SE	PE	ln PE	MM	P(true)	LS
20 Steps, $\alpha = 0.7$	0.60	0.61	0.58	0.66	0.57	0.59	**0.70**

Tab. 5.1.

Analysis of 20-Step Results
Under this specific configuration, Lexical Similarity (LS) achieves the strongest performance with an AUROC of 0.70, followed by length-normalized Predictive Entropy (ln PE) with 0.66. Our proposed Credal Semantic Entropy yields an AUROC of 0.58, which underperforms standard Semantic Entropy (0.61) and equals raw Predictive Entropy (0.58).

5.2.2 Evaluation: 50-Step Interval

To investigate whether a finer resolution of the temperature space improves the calibration of our credal set, we extend our analysis to the 50-step interval (Table A.2). As in the previous experiment, we first evaluate the most restrictive threshold, $\alpha = 0.9$. Note that we were not able to test more α values for the 50-step interval as it was computationally prohibitive.

Table 5.2: UQ performance (AUROC) on the CoQA subset using a 50-step temperature interval and a strict relative likelihood threshold ($\alpha = 0.9$). Metrics compared are our proposed Credal Semantic Entropy (Credal SE), standard Semantic Entropy (SE), Predictive Entropy (PE), length-normalized Predictive Entropy (ln PE), Margin Measure (MM), P(true) and Lexical Similarity (LS).

Configuration	Credal SE	SE	PE	ln PE	MM	P(true)	LS
50 Steps, $\alpha = 0.9$	0.56	0.62	0.55	0.65	0.54	0.62	**0.68**

Tab. 5.2.

Analysis of 50-Step Results
Table 5.2 demonstrates that under the 50-step configuration at $\alpha = 0.9$, Lexical Similarity (LS) maintains its position as the strongest baseline with an AUROC of o.68, again followed by length-normalized Predictive Entropy (ln PE). Standard Semantic Entropy (SE) and P(true) perform competitively at 0.62.

Crucially, our proposed Credal Semantic Entropy yields an AUROC of 0.56. Notably, this is a slight degradation compared to its performance on the coarser 20-step grid (0.58).

5.2.3 Additional Evaluation: Local Analysis at $\tau = 0.5$

In our primary evaluations, the single-point baselines were evaluated at $\tau \approx 1.39$, as this was the temperature that maximized the average likelihood of the calibration data (τ^{ML}). However, Kuhn et al. [KGF23] reported that Semantic Entropy achieves peak performance at a significantly lower temperature, specifically $\tau = 0.5$. To ensure a fair and comprehensive comparison against this baseline, we conducted an additional evaluation centered around this lower temperature regime.

Methodological Caveat: Local Relative Likelihood
Applying our framework strictly, a temperature of $\tau = 0.5$ possesses a global relative likelihood close to zero when compared to the true optimum at $\tau \approx 1.39$, and would thus be excluded by any reasonable α-cut. Consequently, for this specific evaluation, we compute a *local* relative likelihood. We restricted our grid to a narrow 15-step interval around the baseline's optimum, specifically $[0.490, 0.505]$ with a step size of 0.001. Within this local grid, the maximum likelihood was observed at $\tau = 0.505$. Applying an α-threshold of 0.9 yielded a valid temperature set of $\mathcal{T}_{0.9} = \{0.497, 0.498, \ldots, 0.505\}$. It should be noted that the experimental setup for this specific evaluation differs slightly from subsequent experiments, as it was conducted during an earlier phase of the research. While the baseline metrics were evaluated strictly at $\tau = 0.5$, the AUROC for our proposed Credal Semantic Entropy was computed using sequences generated at the local maximum, $\tau^{ML} = 0.505$.

Furthermore, this evaluation utilized a different data partitioning strategy, leveraging approximately 50% of the CoQA dataset for calibration and generation—a notably larger sample size than what was used in later evaluations.

Results

Table 5.3: UQ detection performance (AUROC) on the CoQA subset using a 15-step temperature interval and a strict relative likelihood threshold ($\alpha = 0.9$). Metrics compared are our proposed Credal Semantic Entropy (Credal SE), standard Semantic Entropy (SE), Predictive Entropy (PE), length-normalized Predictive Entropy (ln PE), Margin Measure (MM), P(true) and Lexical Similarity (LS).

Configuration	Credal SE	SE	PE	ln PE	MM	P(true)	LS
15 Steps, $\alpha = 0.9$	0.67	**0.74**	0.64	0.72	0.61	0.69	0.61

Tab. 5.3.

Under this configuration, Semantic Entropy (SE) achieves the highest overall AUROC of 0.74, closely followed by length-normalized Predictive Entropy (0.72). Our proposed Credal Semantic Entropy achieves a score of 0.67, outperforming raw Predictive Entropy (0.64) and Lexical Similarity (0.61), but falling short of the leading baselines.

5.3 Discussion

The empirical evaluation of Credal Semantic Entropy presents a nuanced picture of UQ in LLMs. While the credal framework offers a mathematically rigorous approach to modelling epistemic uncertainty, translating this theoretical advantage into improved UQ proves challenging in practice.

Across nearly all configurations, single-point baselines, particularly length-normalized Predictive Entropy (ln PE) and Semantic Entropy, demonstrated superior performance. In the 20- and 50-step evaluation (5.1), Lexical Similarity achieved the highest AUROC (0.70). This suggests that for the OPT-6.7B model, using a higher temperature, on the CoQA dataset, simple string-level variation across sampled answers is an exceptionally strong proxy for UQ.

As illustrated in Figure 5.1, the performance of Credal Semantic Entropy is highly volatile with respect to the α-cut. The line plot reveals a distinct peak at $\tau = 0.7$ (AUROC = 0.60), with performance degrading when the threshold is made either stricter ($\alpha = 0.9$) or more permissive ($\alpha = 0.6$).

This behaviour aligns with the geometry of credal sets. A highly restrictive threshold ($\alpha = 0.9$) strips the credal set of meaningful distributional diversity, collapsing

the probability bounds into a noisy point estimate. Conversely, a threshold that is too loose ($\alpha = 0.6$) introduces "garbage" temperature-distributions that are poorly calibrated to the data. The peak at $\alpha = 0.7$ captures genuine epistemic uncertainty without admitting uncalibrated noise.

Intuitively, evaluating a denser temperature grid should yield a more accurate approximation of the credal set. However, our comparison between the 20-step and 50-step intervals reveals the opposite: performance degraded from 0.58 to 0.56 (at $\alpha = 0.9$). A plausible explanation for this discrepancy lies in the discrete selection of the maximum likelihood temperature (τ^{ML}). As shown in Section 5.2.3, almost all baseline UQ methods demonstrate superior performance at lower temperatures (e.g. $\tau = 0.5$). Because the coarser 20-step grid is restricted to fewer evaluation points, it isolates $\tau^{ML} = 1.347$, which consequently shifts the boundaries of the resulting α-cut downward. This inadvertently forces the credal set to encompass lower, higher-performing temperatures, thereby inflating the overall evaluation metrics.

This grid-resolution artifact serves as a clear symptom of a much broader, fundamental tension within our methodology: the misalignment between predictive likelihood and UQ. Our framework constructs the credal set around $\tau \approx 1.39$ because this temperature strictly maximizes the likelihood of the calibration data. However, the empirical results show that uncertainty metrics perform much better at UQ at a significantly lower temperature ($\tau = 0.5$, where SE achieves 0.74 AUROC). The temperature that best explains the empirical distribution of the text (maximum likelihood) is not the temperature that separates factual knowledge from high uncertainty. Because our relative likelihood criterion inherently anchors the credal set to the high-temperature, maximum likelihood regime, it forces the evaluation into a parameter space where UQ is more difficult.

Ultimately, while upper semantic entropy over a credal set captures worst-case uncertainty bounds, it effectively acts as a measure of *total* uncertainty. Because it does not explicitly disentangle and subtract the aleatoric component, it removes the advancement over the single-point estimate baselines.

Conclusion 6

6.1 Summary

As LLMs are increasingly deployed in high-stakes domains, their tendency to hallucinate plausible but factually incorrect information remains a critical vulnerability. Standard uncertainty quantification (UQ) methods rely on single, precise probability distributions. By mathematically conflating benign linguistic diversity (aleatoric uncertainty) with a genuine lack of knowledge (epistemic uncertainty), models frequently appear misleadingly confident in their hallucinations.

To address this theoretical limitation, this thesis introduced **Credal Semantic Entropy**, a novel framework grounded in the theory of imprecise probabilities. We hypothesized that true epistemic uncertainty cannot be reliably captured by a single fixed decoding strategy. Instead, we shifted the locus of uncertainty estimation from a single "point-estimate" to a **credal set** - a convex set of plausible probability distributions. Our methodology constructed this set by systematically varying the generation temperature, utilizing the statistical concept of **relative likelihood** to filter out implausible temperatures via an α-cut. We then clustered the generated sequences by semantic meaning using an auxiliary Natural Language Inference model, deriving upper and lower probability bounds for every fundamental meaning.

We empirically evaluated our approach on the conversational Question Answering benchmark, CoQA, using the OPT-6.7B model. Across multiple temperature grids (20-step and 50-step) and varying α thresholds, we compared the upper semantic entropy of our credal set against established single-point baselines.

Our empirical results revealed that it faces significant practical challenges. Under our tested configurations, strong standard baselines outperformed the proposed Credal Semantic Entropy in UQ (measured via AUROC).

In conclusion, this thesis provides a formal theoretical connection between language model decoding hyperparameters and imprecise probabilities. Building upon this foundation, future research can continue to advance interval-based UQ.

6.2 Limitations

While the proposed Credal Semantic Entropy framework provides a rigorous, theoretically grounded approach to UQ, this study is subject to several practical and methodological limitations. The most significant limitation of our approach is the high computational cost incurred during inference. Because of these computational demands and the time constraint of this thesis, our empirical evaluation was restricted in scope. We evaluated the method on a subset of approximately 2400 queries from a single dataset (CoQA [Guo+21]) with only one seed. Evaluating the method across a wider variety of domains and with more granular temperature grids was computationally prohibitive. Furthermore, all experiments were conducted on the OPT-6.7B model [Zha+22]. This specific architecture was selected because it is an open-weights ("white-box") model, granting full access to token-level logits required for our calculations, while also being small enough to run on accessible hardware without incurring the prohibitive API costs associated with proprietary models. However, it remains an open question whether the uncertainty dynamics observed in a 6.7-billion parameter model perfectly mirrors those of massive, instruction-tuned foundation models (e.g. GPT-4 [Ope+24]).

6.3 Future Work

While this thesis establishes a rigorous baseline for Credal Semantic Entropy using temperature scaling, several avenues remain for refining the methodology and extending its application. We outline three primary directions for future research.

6.3.1 Generalization Across Domains and Architecture

Our experiments focused on the CoQA dataset using the OPT-6.7B model. To validate the universality of the credal set approach, future work should assess performance on:

- **diverse benchmarks** such as TriviaQA [Jos+17] and SQuAD [RJL18];
- **safety-critical domains** like healthcare (MedQA [Jin+20]) or legal analysis, where the cost of epistemic failure is significantly higher;
- **state-of-the-art foundation models** such as Gemini [Tea+25] or GPT-4.

6.3.2 Multi-Dimensional Decoding Ensembles

In this work, we constructed the credal set solely by varying the *temperature* parameter. However, the decoding space of modern LLMs is high-dimensional. A natural extension is to define the credal set over a broader range of decoding hyperparameters, including: **Nucleus (Top-p) Sampling** (Varying the cumulative probability threshold [Hol+20]), **Top-k Sampling** (Varying the vocabulary truncation size [FLD18]), and **Beam Search** (comparing stochastic sampling against deterministic beam search paths [FA17]). By defining the α-cut over this joint hyperparameter space, we could capture a more holistic view of the model's inherent uncertainty.

6.3.3 Methodological Extensions: Ensembles and Black-Box Settings

Finally, the relative likelihood framework could be extended beyond single-model decoding. **Model Ensembles:** Instead of perturbing the temperature of a single model, one could construct a credal set using an ensemble of different model checkpoints or architectures, effectively bridging the gap between our approach and the original parameter-space credal prediction [Löh+25]. **Discrete Semantic Entropy:** Currently, our method requires access to token-level log-probabilities (White-Box). To extend this method to closed-source "Black-Box" APIs (e.g., Claude [Ant24], GPT-4 [Ope+24]), future work could integrate **Discrete Semantic Entropy** [Far+24], approximating the cluster probabilities using purely text-based sampling without access to internal logits.

Appendix A

A.1 Experimental Details

Our implementation builds upon the codebase provided by Kuhn et al. [KGF23]. We utilize the HuggingFace Transformers library to implement both the generative model (OPT-6.7B, [Zha+22]) and the entailment model (DeBERTa-large-mnli, [He+21]).

To ensure a fair comparison, we replicate the generation parameters used in the baseline studies. The generation process is divided into two distinct phases using the `generate()` function:

Primary Answer Generation: To produce the model's "best guess" prediction (which is evaluated against the ground truth), we employ beam search with sampling enabled. Specifically, we set the beam size to 5 (`num_beams=5`) and enable multinomial sampling (`do_sample=True`).

Uncertainty Estimation Sampling: To generate the diverse set of sequences required for computing semantic entropy and our proposed credal metrics, we disable beam search (`num_beams=1`) and rely solely on multinomial sampling (`do_sample=True`). Following the findings of Kuhn et al. [KGF23], who demonstrated that uncertainty estimation performance saturates beyond a small number of samples, we generate $N = 10$ stochastic sequences per question.

Detailed implementation details and instructions regarding hardware requirements, environment setup, and the exact commands required to reproduce the tables and figures presented in Section 5 are provided in the repository (`https://github.com/liamhpr/Credal-Prediction-for-LLMs.git`).

The experiments run on one 80GB NVIDIA A100.

For our evaluation, we split the temperature interval $[0.4, 2.2]$ using 20 and 50 steps. The step size stays consistent and the resulting temperature spaces can be seen in the two Tables A.1, A.2 below.

Table A.1: The temperature interval $[0.4, 2.2]$ split into 20 evenly spaced steps. The temperature yielding the maximum average likelihood ($\tau^{ML} = 1.347$) is highlighted in bold. In addition, the resulting α-cuts $\mathcal{T}_\alpha$ with $\alpha \in \{0.6, 0.7, 0.8, 0.9\}$ are shown.

Set of Evaluated Temperatures ($\mathcal{T}$)	
0.400, 0.495, 0.590, 0.684, 0.779, 0.874, 0.968, 1.063, 1.158, 1.253, **1.347**, 1.442, 1.537, 1.632, 1.726, 1.821, 1.916, 2.011, 2.105, 2.200	
$\mathcal{T}_{0.9} = \{1.253, \ldots, 1.537\}$	$\lvert\mathcal{T}_{0.9}\rvert = 4$
$\mathcal{T}_{0.8} = \{1.158, \ldots, 1.632\}$	$\lvert\mathcal{T}_{0.8}\rvert = 6$
$\mathcal{T}_{0.7} = \{1.063, \ldots, 1.726\}$	$\lvert\mathcal{T}_{0.7}\rvert = 10$
$\mathcal{T}_{0.6} = \{1.063, \ldots, 1.821\}$	$\lvert\mathcal{T}_{0.6}\rvert = 11$

Tab. A.1.

Table A.2: The temperature interval $[0.4, 2.2]$ split into 50 evenly spaced steps. The temperature yielding the maximum average likelihood ($\tau^{ML} = 1.392$) is highlighted in bold. In addition, the resulting α-cuts $\mathcal{T}_\alpha$ with $\alpha = 0.9$ are shown.

Set of Evaluated Temperatures ($\mathcal{T}$)	
0.400, 0.437, 0.474, 0.510, 0.547, 0.584, 0.620, 0.657, 0.694, 0.731, 0.767, 0.804, 0.841, 0.878, 0.914, 0.951, 0.988, 1.025, 1.061, 1.098, 1.135, 1.171, 1.208, 1.245, 1.282, 1.318, 1.355, **1.392**, 1.429, 1.465, 1.502, 1.539, 1.576, 1.612, 1.649, 1.686, 1.723, 1.759, 1.796, 1.833, 1.869, 1.906, 1.943, 1.980, 2.016, 2.053, 2.090, 2.127, 2.163, 2.200	
$\mathcal{T}_{0.9} = \{1.208, \ldots, 1.576\}$	$\lvert\mathcal{T}_{0.9}\rvert = 11$

Tab. A.2.

Bibliography

[AKM06] J. Abellán, G.J. Klir, and S. Moral. "Disaggregated total uncertainty measure for credal sets". In: *International Journal of General Systems* 35.1 (2006), pp. 29–44. eprint: https://doi.org/10.1080/03081070500473490 (cit. on p. 8).

[AM05] Joaquín Abellán and Serafín Moral. "Upper entropy of credal sets. Applications to credal classification". In: *International Journal of Approximate Reasoning* 39.2 (2005). Imprecise Probabilities and Their Applications, pp. 235–255 (cit. on pp. 8, 9).

[Ant24] Anthropic. *Introducing the next generation of Claude*. 2024 (cit. on p. 39).

[ACC12] Alessandro Antonucci, Marco E. G. V. Cattaneo, and Giorgio Corani. "Likelihood-Based Robust Classification with Bayesian Networks". In: *Advances in Computational Intelligence. 14th International Conference on Information Processing and Management of Uncertainty in Knowledge-Based Systems, IPMU 2012, Catania, Italy, July 9-13, 2012, Proceedings, Part III*. Ed. by Salvatore Greco, Bernadette Bouchon-Meunier, Giulianella Coletti, et al. Berlin: Springer, 2012, pp. 491–500 (cit. on p. 14).

[BCB16] Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. *Neural Machine Translation by Jointly Learning to Align and Translate*. 2016. arXiv: 1409.0473 [cs.CL] (cit. on p. 9).

[Bed+25] Suhana Bedi, Yutong Liu, Lucy Orr-Ewing, et al. "Testing and Evaluation of Health Care Applications of Large Language Models: A Systematic Review". In: *JAMA* 333.4 (Jan. 2025), pp. 319–328. eprint: https://jamanetwork.com/journals/jama/articlepdf/2825147/jama_bedi_2024_oi_240124_1737130172.28802.pdf (cit. on pp. 1, 5).

[Bir62] Allan Birnbaum. "On the Foundations of Statistical Inference". In: *Journal of the American Statistical Association* 57.298 (1962), pp. 269–306. eprint: https://doi.org/10.1080/01621459.1962.10480660 (cit. on p. 13).

[Bom+22] Rishi Bommasani, Drew A. Hudson, Ehsan Adeli, et al. *On the Opportunities and Risks of Foundation Models*. 2022. arXiv: 2108.07258 [cs.LG] (cit. on p. 1).

[Bra97] Andrew P. Bradley. "The use of the area under the ROC curve in the evaluation of machine learning algorithms". In: *Pattern Recognition* 30.7 (1997), pp. 1145–1159 (cit. on p. 31).

[Bri90] John S. Bridle. “Probabilistic Interpretation of Feedforward Classification Network Outputs, with Relationships to Statistical Pattern Recognition”. In: *Neurocomputing*. Ed. by Françoise Fogelman Soulié and Jeanny Hérault. Berlin, Heidelberg: Springer Berlin Heidelberg, 1990, pp. 227–236 (cit. on p. 9).

[Bro+20] Tom B. Brown, Benjamin Mann, Nick Ryder, et al. *Language Models are Few-Shot Learners*. 2020. arXiv: 2005.14165 [cs.CL] (cit. on pp. 1, 29).

[CHM11] Luis de Campos, Juan Huete, and Serafín Moral. “Probability interval; A tool for uncertain reasoning”. In: *International Journal of Uncertainty, Fuzziness and Knowledge-Based Systems* 02 (Nov. 2011) (cit. on p. 8).

[CH79] D.R. Cox and D.V. Hinkley. *Theoretical Statistics*. Taylor & Francis, 1979 (cit. on p. 13).

[Coz00] Fabio G. Cozman. “Credal networks”. In: *Artificial Intelligence* 120.2 (2000), pp. 199–233 (cit. on p. 7).

[DD20] Shrey Desai and Greg Durrett. *Calibration of Pre-trained Transformers*. 2020. arXiv: 2003.07892 [cs.CL] (cit. on p. 1).

[FLD18] Angela Fan, Mike Lewis, and Yann Dauphin. *Hierarchical Neural Story Generation*. 2018. arXiv: 1805.04833 [cs.CL] (cit. on pp. 11, 39).

[Far+24] Sebastian Farquhar, Jannik Kossen, Lorenz Kuhn, and Yarin Gal. “Detecting hallucinations in large language models using semantic entropy”. In: *Nature* 630.8017 (2024), pp. 625–630 (cit. on p. 39).

[Fom+20] Marina Fomicheva, Shuo Sun, Lisa Yankovskaya, et al. *Unsupervised Quality Estimation for Neural Machine Translation*. 2020. arXiv: 2005.10608 [cs.CL] (cit. on p. 1).

[FA17] Markus Freitag and Yaser Al-Onaizan. “Beam Search Strategies for Neural Machine Translation”. In: *Proceedings of the First Workshop on Neural Machine Translation*. Ed. by Thang Luong, Alexandra Birch, Graham Neubig, and Andrew Finch. Vancouver: Association for Computational Linguistics, Aug. 2017, pp. 56–60 (cit. on pp. 10, 39).

[GG16] Yarin Gal and Zoubin Ghahramani. *Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning*. 2016. arXiv: 1506.02142 [stat.ML] (cit. on p. 15).

[GR07] Tilmann Gneiting and Adrian E Raftery. “Strictly Proper Scoring Rules, Prediction, and Estimation”. In: *Journal of the American Statistical Association* 102.477 (2007), pp. 359–378. eprint: https://doi.org/10.1198/016214506000001437 (cit. on p. 6).

[GBC16] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. http://www.deeplearningbook.org. MIT Press, 2016 (cit. on p. 13).

[Guo+17] Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. *On Calibration of Modern Neural Networks*. 2017. arXiv: 1706.04599 [cs.LG] (cit. on pp. 1, 5–7).

[Guo+21] Meiqi Guo, Mingda Zhang, Siva Reddy, and Malihe Alikhani. "Abg-CoQA: Clarifying Ambiguity in Conversational Question Answering". In: *3rd Conference on Automated Knowledge Base Construction*. 2021 (cit. on p. 38).

[He+21] Pengcheng He, Xiaodong Liu, Jianfeng Gao, and Weizhu Chen. *DeBERTa: Decoding-enhanced BERT with Disentangled Attention*. 2021. arXiv: 2006.03654 [cs.CL] (cit. on pp. 17, 24, 29, 41).

[HG18] Dan Hendrycks and Kevin Gimpel. *A Baseline for Detecting Misclassified and Out-of-Distribution Examples in Neural Networks*. 2018. arXiv: 1610.02136 [cs.NE] (cit. on p. 6).

[HVD+15] Geoffrey Hinton, Oriol Vinyals, Jeff Dean, et al. "Distilling the knowledge in a neural network". In: *arXiv preprint arXiv:1503.02531* (2015) (cit. on pp. 11, 12).

[Hol+20] Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. *The Curious Case of Neural Text Degeneration*. 2020. arXiv: 1904.09751 [cs.CL] (cit. on pp. 10, 11, 39).

[Hua+23] Lei Huang, Weijiang Yu, Weitao Ma, et al. "A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions". In: *ACM Transactions on Information Systems* 43 (2023), pp. 1–55 (cit. on p. 1).

[HW21] Eyke Hüllermeier and Willem Waegeman. "Aleatoric and epistemic uncertainty in machine learning: an introduction to concepts and methods". In: *Machine Learning* 110.3 (Mar. 2021), pp. 457–506 (cit. on pp. 1, 5, 7, 8, 12).

[Ji+23] Ziwei Ji, Nayeon Lee, Rita Frieske, et al. "Survey of Hallucination in Natural Language Generation". In: *ACM Computing Surveys* 55.12 (Mar. 2023), pp. 1–38 (cit. on p. 1).

[Jin+20] Di Jin, Eileen Pan, Nassim Oufattole, et al. *What Disease does this Patient Have? A Large-scale Open Domain Question Answering Dataset from Medical Exams*. 2020. arXiv: 2009.13081 [cs.CL] (cit. on p. 38).

[Jos+17] Mandar Joshi, Eunsol Choi, Daniel S. Weld, and Luke Zettlemoyer. *TriviaQA: A Large Scale Distantly Supervised Challenge Dataset for Reading Comprehension*. 2017. arXiv: 1705.03551 [cs.CL] (cit. on pp. 25, 38).

[Kad+22] Saurav Kadavath, Tom Conerly, Amanda Askell, et al. *Language Models (Mostly) Know What They Know*. 2022. arXiv: 2207.05221 [cs.CL] (cit. on pp. 16, 30).

[KG17] Alex Kendall and Yarin Gal. *What Uncertainties Do We Need in Bayesian Deep Learning for Computer Vision?* 2017. araXiv: 1703.04977 (cs.CV) (cit. on pp. 5, 12).

[KGF23] Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. *Semantic Uncertainty: Linguistic Invariances for Uncertainty Estimation in Natural Language Generation*. 2023. arXiv: 2302.09664 [cs.CL] (cit. on pp. 1, 2, 12, 16–18, 29, 30, 33, 41).

[LPB17] Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell. *Simple and Scalable Predictive Uncertainty Estimation using Deep Ensembles*. 2017. arXiv: 1612.01474 [stat.ML] (cit. on p. 15).

[Lev80] Isaac Levi. *The Enterprise of Knowledge: An Essay on Knowledge, Credal Probability, and Chance*. MIT Press, 1980 (cit. on p. 7).

[Lin04] Chin-Yew Lin. "ROUGE: A Package for Automatic Evaluation of Summaries". In: *Text Summarization Branches Out*. Barcelona, Spain: Association for Computational Linguistics, July 2004, pp. 74–81 (cit. on pp. 16, 30).

[LLS22] Zi Lin, Jeremiah Zhe Liu, and Jingbo Shang. "Towards Collaborative Neural-Symbolic Graph Semantic Parsing via Uncertainty". In: *Findings of the Association for Computational Linguistics: ACL 2022*. Ed. by Smaranda Muresan, Preslav Nakov, and Aline Villavicencio. Dublin, Ireland: Association for Computational Linguistics, May 2022, pp. 4160–4173 (cit. on pp. 15, 30).

[Löh+25] Timo Löhr, Paul Hofman, Felix Mohr, and Eyke Hüllermeier. *Credal Prediction based on Relative Likelihood*. 2025. arXiv: 2505.22332 [stat.ML] (cit. on pp. 2, 13, 19, 23, 39).

[MG21] Andrey Malinin and Mark Gales. *Uncertainty Estimation in Autoregressive Structured Prediction*. 2021. arXiv: 2002.07650 [stat.ML] (cit. on pp. 1, 15, 25, 30).

[MC18] Kenton Murray and David Chiang. *Correcting Length Bias in Neural Machine Translation*. 2018. arXiv: 1808.10006 [cs.CL] (cit. on p. 13).

[Nav+24] Humza Naveed, Asad Ullah Khan, Shi Qiu, et al. *A Comprehensive Overview of Large Language Models*. 2024. arXiv: 2307.06435 [cs.CL] (cit. on p. 9).

[Ope+24] OpenAI, Josh Achiam, Steven Adler, et al. *GPT-4 Technical Report*. 2024. arXiv: 2303.08774 [cs.CL] (cit. on pp. 38, 39).

[Pap+02] Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. "BLEU: a method for automatic evaluation of machine translation". In: *Proceedings of the 40th Annual Meeting on Association for Computational Linguistics*. ACL '02. Philadelphia, Pennsylvania: Association for Computational Linguistics, 2002, pp. 311–318 (cit. on p. 16).

[Pin88] R. Pincus. "Aitchison, J.: The Statistical Analysis of Compositional Data. Chapman and Hall, London - New York 1986, XII, 416 pp., £ 25,00". In: *Biometrical Journal* 30.7 (1988), pp. 794–794. eprint: https://onlinelibrary.wiley.com/doi/pdf/10.1002/bimj.4710300705 (cit. on p. 7).

[RJL18] Pranav Rajpurkar, Robin Jia, and Percy Liang. *Know What You Don't Know: Unanswerable Questions for SQuAD*. 2018. arXiv: 1806.03822 [cs.CL] (cit. on p. 38).

[RCM19] Siva Reddy, Danqi Chen, and Christopher D. Manning. *CoQA: A Conversational Question Answering Challenge*. 2019. arXiv: 1808.07042 [cs.CL] (cit. on pp. 1, 29).

[Sha48] C. E. Shannon. “A Mathematical Theory of Communication”. In: *Bell System Technical Journal* 27.3 (1948), pp. 379–423. eprint: https://onlinelibrary.wiley.com/doi/pdf/10.1002/j.1538-7305.1948.tb01338.x (cit. on p. 6).

[Sii+25] Marco Siino, Mariana Falco, Daniele Croce, and Paolo Rosso. “Exploring LLMs Applications in Law: A Literature Review on Current Legal NLP Approaches”. In: *IEEE Access* 13 (2025), pp. 18253–18276 (cit. on p. 1).

[SVL14] Ilya Sutskever, Oriol Vinyals, and Quoc V. Le. *Sequence to Sequence Learning with Neural Networks*. 2014. arXiv: 1409.3215 [cs.CL] (cit. on p. 10).

[Tea+25] Gemini Team, Rohan Anil, Sebastian Borgeaud, et al. *Gemini: A Family of Highly Capable Multimodal Models*. 2025. arXiv: 2312.11805 [cs.CL] (cit. on p. 38).

[Tia+23] Haoye Tian, Weiqi Lu, Tsz On Li, et al. *Is ChatGPT the Ultimate Programming Assistant – How far is it?* 2023. arXiv: 2304.11938 [cs.SE] (cit. on p. 9).

[Vas+17] Ashish Vaswani, Noam Shazeer, Niki Parmar, et al. “Attention is All you Need”. In: *Advances in Neural Information Processing Systems*. Ed. by I. Guyon, U. Von Luxburg, S. Bengio, et al. Vol. 30. Curran Associates, Inc., 2017 (cit. on pp. 9, 10).

[WM99] P. Walley and S. Moral. “Upper probabilities based only on the likelihood function”. In: *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 61.4 (1999), pp. 831–847. eprint: https://rss.onlinelibrary.wiley.com/doi/pdf/10.1111/1467-9868.00205 (cit. on p. 13).

[Wal91] Peter Walley. *Statistical Reasoning with Imprecise Probabilities*. Chapman & Hall, 1991 (cit. on p. 7).

[Was90] Larry A. Wasserman. “Belief functions and statistical inference”. In: *Canadian Journal of Statistics* 18.3 (1990), pp. 183–196. eprint: https://onlinelibrary.wiley.com/doi/pdf/10.2307/3315449 (cit. on p. 13).

[Wei+23] Jason Wei, Xuezhi Wang, Dale Schuurmans, et al. *Chain-of-Thought Prompting Elicits Reasoning in Large Language Models*. 2023. arXiv: 2201.11903 [cs.CL] (cit. on pp. 1, 9).

[Wil38] S. S. Wilks. “The Large-Sample Distribution of the Likelihood Ratio for Testing Composite Hypotheses”. In: *The Annals of Mathematical Statistics* 9.1 (1938), pp. 60–62 (cit. on p. 13).

[WR16] Sam Wiseman and Alexander M. Rush. “Sequence-to-Sequence Learning as Beam-Search Optimization”. In: *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*. Ed. by Jian Su, Kevin Duh, and Xavier Carreras. Austin, Texas: Association for Computational Linguistics, Nov. 2016, pp. 1296–1306 (cit. on p. 10).

[Zha+22] Susan Zhang, Stephen Roller, Naman Goyal, et al. *OPT: Open Pre-trained Transformer Language Models*. 2022. arXiv: 2205.01068 [cs.CL] (cit. on pp. 29, 38, 41).

List of Figures

3.1. Visualization of semantic clustering for the question "What is the capital of Germany?". After the generation of seven answer strings, they are semantically clustered and afterwards the the cluster probabilities are computed as shown in 3.1. 17

4.1. Visualization of the first three stages of our methodology. The process begins with multi-temperature generation, followed by semantic clustering performed on the union of all generated sequences. Finally, the probability mass of each cluster is marginalized at every temperature to establish the probability intervals used for UQ. 22

4.2. Geometric representation of the credal set derived from the probability intervals established in Figure 4.1. Each axis of the radar chart corresponds to a distinct semantic cluster. The blue bars represent the probability intervals (lower to upper bounds) for each class, illustrating the model's epistemic uncertainty across the potential predictions. This visualization approach follows [Hof+26]. 26

5.1. Impact of the relative likelihood threshold (α) on the UQ performance of Credal Semantic Entropy. The line plot illustrates the AUROC scores evaluated on the 20-step temperature grid for $\alpha \in \{0.6, 0.7, 0.8, 0.9\}$. . . 32

List of Tables

5.1. 32
5.2. 33
5.3. 34

A.1. 42
A.2. 42

Ich erkläre hiermit, dass ich die vorliegende Arbeit selbstständig und nur mit den angegebenen Hilfsmitteln verfasst habe. Alle Passagen, die ich aus einer Literatur oder anderen Quellen übernommen habe, sind eindeutig als Zitate mit Quellenangabe gekennzeichnet. Diese gedruckte Fassung ist ein Ausdruck der eingereichten elektronischen Fassung.

I hereby declare that I have written this work independently and only with the aids specified. All passages that I have taken from literature or other sources have been clearly marked as quotations with reference to the source. This printed copy is a printout of the submitted electronic copy.

_____

Datum/date Unterschrift/signature

Colophon

This thesis was typeset with LaTeX 2_ε. It uses the *Clean Thesis* style developed by Ricardo Langner. The design of the *Clean Thesis* style is inspired by user guide documents from Apple Inc.

Download the *Clean Thesis* style at `http://cleanthesis.der-ric.de/`.